

# AJLA Concierge API - Deployment Guide

---

## Overview

This guide covers deploying the AJLA Concierge API to various cloud platforms. The application is a FastAPI backend with PostgreSQL database and WebSocket support.

---

## Prerequisites

Before deploying, ensure you have:

- Git repository with latest code pushed
- PostgreSQL database (provided by platform or external service)
- Environment variables ready:
  - `DATABASE_URL` - PostgreSQL connection string
  - `JWT_SECRET_KEY` - Random secure string for token signing

## Generate JWT Secret Key

```
python3 -c "import secrets; print(secrets.token_urlsafe(32))"
```

---

## Option 4: AWS Deployment

AWS offers multiple deployment options. Here are two approaches:

### Option 4A: AWS App Runner (Easiest)

App Runner is AWS's fully managed container service - similar to Railway/Render.

#### Step 1: Prerequisites

- AWS Account
- AWS CLI installed and configured
- Docker image pushed to ECR (or connect GitHub directly)

## Step 2: Create RDS PostgreSQL Database

1. Go to **AWS Console → RDS**
2. Click "**Create database**"
3. Configure:
  - o **Engine:** PostgreSQL 16
  - o **Template:** Free tier
  - o **DB instance identifier:** `ajla-db`
  - o **Master username:** `ajla_user`
  - o **Master password:** (save this!)
  - o **Public access:** Yes (for initial setup)
  - o **VPC security group:** Create new, allow inbound on port 5432
4. Click "**Create database**"
5. Wait for status to become "Available"
6. Copy the **Endpoint** (e.g., `ajla-db.xxxxxx.us-east-1.rds.amazonaws.com`)

Your DATABASE\_URL will be:

```
postgresql://ajla_user:YOUR_PASSWORD@ajla-db.xxxxxx.us-east-  
1.rds.amazonaws.com:5432/postgres
```

## Step 3: Create App Runner Service

1. Go to **AWS Console → App Runner**
2. Click "**Create service**"
3. **Source:**
  - o Select "**Source code repository**"
  - o Connect to GitHub and select your repo
  - o Branch: `main`
4. **Build settings:**
  - o **Runtime:** Python 3.9
  - o **Build command:** `pip install -r requirements.txt`
  - o **Start command:** `uvicorn main:app --host 0.0.0.0 --port 8080`
  - o **Port:** 8080
5. **Service settings:**
  - o **Service name:** `ajla-api`
  - o **CPU:** 0.25 vCPU
  - o **Memory:** 0.5 GB
6. **Environment variables:**

Variable	Value
----------	-------

Variable	Value
DATABASE_URL	postgresql://ajla user:PASS@ajla-db.xxx.rds.amazonaws.com:5432/postgres
JWT_SECRET_KEY	your-generated-secret-key
JWT_ALGORITHM	HS256
DEBUG	false
PORT	8080

7. Click "**Create & deploy**"

8. Your API will be at: <https://xxxxx.us-east-1.awsapprunner.com>

---

## Option 4B: AWS EC2 + Docker (Full Control)

For more control and cost optimization on a VPS.

### Step 1: Launch EC2 Instance

1. Go to **AWS Console** → **EC2** → **Launch Instance**

2. Configure:

- o **Name:** ajla-api-server
- o **AMI:** Ubuntu 22.04 LTS
- o **Instance type:** t3.micro (free tier eligible)
- o **Key pair:** Create new or use existing
- o **Security group:** Allow ports 22 (SSH), 80 (HTTP), 443 (HTTPS), 8000 (API)

3. Click "**Launch instance**"

### Step 2: Connect to EC2

```
ssh -i your-key.pem ubuntu@your-ec2-public-ip
```

### Step 3: Install Docker

```
# Update system
sudo apt update && sudo apt upgrade -y

# Install Docker
curl -fsSL https://get.docker.com | sh
```

```
sudo usermod -aG docker ubuntu

# Install Docker Compose
sudo apt install docker-compose -y

# Logout and login again for group changes
exit
```

## Step 4: Clone Repository

```
ssh -i your-key.pem ubuntu@your-ec2-public-ip

git clone https://github.com/thebahmed00/conceirgeKsa.git
cd conceirgeKsa
```

## Step 5: Create Production Environment File

```
cat > .env.prod << EOF
DATABASE_URL=postgresql://ajla_user:YOUR_PASSWORD@your-rds-endpoint:5432/
JWT_SECRET_KEY=your-generated-secret-key
JWT_ALGORITHM=HS256
JWT_EXPIRATION_HOURS=24
DEBUG=false
LOG_LEVEL=INFO
PORT=8000
EOF
```

## Step 6: Create Production Docker Compose

```
cat > docker-compose.prod.yml << 'EOF'
version: '3.8'

services:
  api:
    build: .
    ports:
      - "8000:8000"
    env_file:
      - .env.prod
    restart: unless-stopped
    healthcheck:
```

```
test: ["CMD", "curl", "-f", "http://localhost:8000/health"]
interval: 30s
timeout: 10s
retries: 3
EOF
```

## Step 7: Deploy

```
docker-compose -f docker-compose.prod.yml up -d --build
```

## Step 8: Setup Nginx Reverse Proxy (Optional but Recommended)

```
sudo apt install nginx -y
```

```
sudo cat > /etc/nginx/sites-available/ajla << 'EOF'
server {
    listen 80;
    server_name your-domain.com;

    location / {
        proxy_pass http://localhost:8000;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection "upgrade";
        proxy_set_header Host $host;
        proxy_set_header X-Real-IP $remote_addr;
        proxy_set_header X-Forwarded-For $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_read_timeout 86400;  # For WebSocket
    }
}
EOF

sudo ln -s /etc/nginx/sites-available/ajla /etc/nginx/sites-enabled/
sudo nginx -t
sudo systemctl restart nginx
```

## Step 9: Setup SSL with Let's Encrypt (Production)

```
sudo apt install certbot python3-certbot-nginx -y
sudo certbot --nginx -d your-domain.com
```

---

## Option 4C: AWS ECS Fargate (Production-Grade)

For production workloads with auto-scaling.

### Step 1: Push Docker Image to ECR

```
# Create ECR repository
aws ecr create-repository --repository-name ajla-api --region us-east-1

# Login to ECR
aws ecr get-login-password --region us-east-1 | docker login --username F
# Build and push
docker build -t ajla-api .
docker tag ajla-api:latest YOUR_ACCOUNT_ID.dkr.ecr.us-east-1.amazonaws.com/ajla-api:latest
docker push YOUR_ACCOUNT_ID.dkr.ecr.us-east-1.amazonaws.com/ajla-api:latest
```

### Step 2: Create ECS Cluster

1. Go to **AWS Console** → **ECS** → **Create cluster**
2. **Cluster name:** ajla-cluster
3. **Infrastructure:** AWS Fargate
4. Click "**Create**"

### Step 3: Create Task Definition

1. Go to **Task definitions** → **Create new task definition**
2. Configure:
  - o **Family:** ajla-api-task
  - o **Launch type:** Fargate
  - o **CPU:** 0.25 vCPU
  - o **Memory:** 0.5 GB
3. **Container:**
  - o **Name:** ajla-api
  - o **Image URI:** YOUR\_ACCOUNT\_ID.dkr.ecr.us-east-1.amazonaws.com/ajla-api:latest
  - o **Port mappings:** 8000
  - o **Environment variables:** Add all required vars
4. Click "**Create**"

## Step 4: Create Service

1. Go to your cluster → **Services** → **Create**
2. Configure:
  - **Launch type:** Fargate
  - **Task definition:** `ajla-api-task`
  - **Service name:** `ajla-api-service`
  - **Desired tasks:** 2
  - **Networking:** Select VPC, subnets, security group
3. **Load balancer:** Application Load Balancer
4. Click "**Create**"

## Step 5: Configure ALB Health Check

- **Health check path:** `/health`
- **Healthy threshold:** 2
- **Interval:** 30 seconds

---

## AWS Cost Estimates (Monthly)

Service	Free Tier	After Free Tier
<b>App Runner</b>	-	~\$5-15/month
<b>EC2 t3.micro</b>	750 hrs/month (1 year)	~\$8/month
<b>RDS db.t3.micro</b>	750 hrs/month (1 year)	~\$15/month
<b>ECS Fargate</b>	-	~\$10-30/month

---

## AWS Security Best Practices

1. **Never commit credentials** - Use environment variables or AWS Secrets Manager
2. **Use IAM roles** - Don't use root account
3. **Enable RDS encryption** - Encrypt data at rest
4. **VPC security groups** - Restrict database access to only your app
5. **Enable CloudWatch** - Monitor logs and set up alerts
6. **Use Secrets Manager** for sensitive values:

```
# Store secret
aws secretsmanager create-secret --name ajla/jwt-secret --secret-string "
```

```
# Reference in ECS task definition
# Use "valueFrom": "arn:aws:secretsmanager:region:account:secret:ajla/jwt
```

---

## Option 2: Docker (Self-Hosted / VPS)

For deployment on your own server (DigitalOcean, AWS EC2, etc.)

### Step 1: Install Docker on Server

```
curl -fsSL https://get.docker.com | sh
```

### Step 2: Create docker-compose.prod.yml

```
version: '3.8'

services:
  api:
    build: .
    ports:
      - "8000:8000"
    environment:
      - DATABASE_URL=postgresql://ajla_user:your_password@db:5432/ajla_db
      - JWT_SECRET_KEY=your-generated-secret-key
      - JWT_ALGORITHM=HS256
      - DEBUG=false
      - PORT=8000
    depends_on:
      - db
    restart: unless-stopped

  db:
    image: postgres:16-alpine
    environment:
      - POSTGRES_USER=ajla_user
      - POSTGRES_PASSWORD=your_password
      - POSTGRES_DB=ajla_db
    volumes:
      - postgres_data:/var/lib/postgresql/data
    restart: unless-stopped
```

```
volumes:
```

```
    postgres_data:
```

## Step 3: Deploy

```
docker-compose -f docker-compose.prod.yml up -d
```

## Step 4: Setup Reverse Proxy (Nginx)

```
server {  
    listen 80;  
    server_name your-domain.com;  
  
    location / {  
        proxy_pass http://localhost:8000;  
        proxy_http_version 1.1;  
        proxy_set_header Upgrade $http_upgrade;  
        proxy_set_header Connection "upgrade";  
        proxy_set_header Host $host;  
        proxy_set_header X-Real-IP $remote_addr;  
    }  
}
```

## Environment Variables Reference

Variable	Required	Default	Description
DATABASE_URL	<input checked="" type="checkbox"/> Yes	-	PostgreSQL connection string
JWT_SECRET_KEY	<input checked="" type="checkbox"/> Yes	-	Secret key for JWT signing
JWT_ALGORITHM	No	HS256	JWT signing algorithm
JWT_EXPIRATION_HOURS	No	24	Token expiration time
DEBUG	No	true	Enable debug mode
LOG_LEVEL	No	INFO	Logging level
PORT	No	8000	Server port (auto-set by Railway/Render)

# Post-Deployment Checklist

- Health endpoint returns `{"status": "healthy"}`
- API docs accessible at `/docs`
- User registration works (`POST /api/v1/auth/register`)
- User login works (`POST /api/v1/auth/login`)
- WebSocket chat connects (`ws://your-domain/ws/chat/{id}?token=...`)
- Database migrations applied
- Admin user created (if needed)

## Create Admin User

After deployment, create an admin user via SQL:

```
-- First register a user via API, then promote to admin:  
UPDATE users SET is_admin = TRUE WHERE email = 'admin@example.com';
```

---

## Troubleshooting

### 502 Bad Gateway

- Check if the app is binding to `$PORT` environment variable
- Verify health endpoint returns 200
- Check logs for startup errors

### Database Connection Failed

- Verify `DATABASE_URL` format: `postgresql://user:pass@host:port/dbname`
- Check if database is accessible from the app
- For Railway: ensure PostgreSQL service is linked

### JWT Errors

- Ensure `JWT_SECRET_KEY` is set
- Check token format in Authorization header: `Bearer <token>`

### WebSocket Connection Failed

- WebSocket URLs use `wss://` for HTTPS domains
  - Include token as query parameter: `?token=your-jwt-token`
  - Check CORS settings if connecting from browser
- 

## API Endpoints Summary

Method	Endpoint	Description
GET	/	Welcome message
GET	/health	Health check
GET	/docs	Swagger UI
POST	/api/v1/auth/register	User registration
POST	/api/v1/auth/login	User login
GET	/api/v1/auth/me	Current user info
GET	/api/v1/users/me	User profile
PUT	/api/v1/users/me	Update profile
POST	/api/v1/requests	Create request
GET	/api/v1/requests	List requests
GET	/api/v1/conversations	List conversations
GET	/api/v1/conversations/{id}	Get conversation
POST	/api/v1/conversations/{id}/messages	Send message
WS	/ws/chat/{conversation_id}	WebSocket chat
GET	/api/v1/admin/conversations	Admin: List all
POST	/api/v1/admin/conversations/{id}/messages	Admin: Reply

---

## Support

For issues, check:

1. Deploy logs in your platform's dashboard
2. Application logs for error messages
3. Database connectivity
4. Environment variables configuration

